

# SEMT20003, Part 4: Neural Networks

Laurence Aitchison

Most of the content this week is in the Colab notebooks. But there's a few pages here introducing neural networks.

So far, we have worked with linear functions, of the form,


$$\mathbf{f}(\mathbf{x}) = \mathbf{x}\mathbf{W}. \quad (1)$$

And we used this function as the prediction in linear regression (with a squared error loss function), or as the logits in classification (with a maximum-likelihood objective / cross-entropy loss function).


But we can use *any* function here. So what function should we use? Well, there's a few requirements.

- lots of parameters (so gradient descent has lots of knobs to twiddle to improve the function).
- doesn't reduce to something simpler.
- fast (e.g. uses matmuls).

We can draw the current setup as,



Now, one way to get a more interesting function is to stack multiple layers (i.e. connect the outputs of one linear layer to the inputs of another layer):



Note that we have included an “nonlinearity”,  $\phi$ . This nonlinearity is necessary, because without it (or if we set  $\phi(\mathbf{h}) = \mathbf{h}$ ), we end up with,

$$\mathbf{h}_1 = \mathbf{x}\mathbf{W}_1 \quad (2a)$$

$$\mathbf{h}_2 = \mathbf{h}_1\mathbf{W}_2 \quad (2b)$$

$$\mathbf{f} = \mathbf{h}_2\mathbf{W}_3 \quad (2c)$$

If we substitute  $\mathbf{h}_1$  into the expression for  $\mathbf{h}_2$ , and  $\mathbf{h}_2$  into the expression for  $\mathbf{f}$ , we end up with,

$$\mathbf{f} = \mathbf{x} \underbrace{\mathbf{W}_1\mathbf{W}_2\mathbf{W}_3}_{=\mathbf{W}} = \mathbf{x}\mathbf{W}. \quad (3)$$

This is again just a linear combination of the inputs,  $\mathbf{x}$ , so offers no advantages over the original linear model.

Including a nonlinearity prevents the collapse back to a linear model, giving us a strictly more powerful class of models.

$$\mathbf{h}_1 = \phi(\mathbf{x}\mathbf{W}_1) \quad (4a)$$


$$\mathbf{h}_2 = \phi(\mathbf{h}_1\mathbf{W}_2) \quad (4b)$$

$$\mathbf{f} = \mathbf{h}_2\mathbf{W}_3 \quad (4c)$$

There are lots of different choices for the nonlinearity. Perhaps the most common is the “rectified linear unit” or ReLU,

$$\phi_{\text{ReLU}}(a) = \begin{cases} a & \text{if } a > 0 \\ 0 & \text{otherwise} \end{cases} \quad (5)$$


This is commonly known as a 3-layer network, as there are three weight-matrices.



**Q: Why don't we have a nonlinearity in the final layer (i.e. for producing the output,  $\mathbf{f}$ )?** Nonlinearities can impose constraints on their outputs. For instance, ReLU can't return negative numbers. Alternatively, in the 1990's people sometimes used a sigmoid as a nonlinearity, but that can only output numbers between 0 and 1. We don't want to impose such constraints on our outputs, which is why we don't apply the output at the final layer.

**Q: Is there a better way of understanding what nonlinearities are doing than just "not being linear"?** Yes ... but you need quite a bit of theory to get there.

**Q: Why is relu far more commonly used than e.g. sigmoid?** This is to do with propagation of the gradients. Specifically, the gradients of the sigmoid are often very small: if  $h = \phi_{\text{sigmoid}}(a) = \sigma(a)$ ,



then the gradient,  $\partial h / \partial a$ , is small for large positive or large negative  $a$ . That makes gradient descent difficult, as lots of the gradients get too small. In contrast, the gradients for relu don't explode/vanish, as the gradient,  $\partial h / \partial a = 1$  for all  $a > 0$ .

## 1 Exercises

There are no exercises this week! (See CoLab notebooks).